

RANDOM LASSO

BY SIJIAN WANG, BIN NAN¹, SAHARON ROSSET² AND JI ZHU³

*University of Wisconsin, University of Michigan, Tel Aviv University and
University of Michigan*

We propose a computationally intensive method, the random lasso method, for variable selection in linear models. The method consists of two major steps. In step 1, the lasso method is applied to many bootstrap samples, each using a set of randomly selected covariates. A measure of importance is yielded from this step for each covariate. In step 2, a similar procedure to the first step is implemented with the exception that for each bootstrap sample, a subset of covariates is randomly selected with unequal selection probabilities determined by the covariates' importance. Adaptive lasso may be used in the second step with weights determined by the importance measures. The final set of covariates and their coefficients are determined by averaging bootstrap results obtained from step 2. The proposed method alleviates some of the limitations of lasso, elastic-net and related methods noted especially in the context of microarray data analysis: it tends to remove highly correlated variables altogether or select them all, and maintains maximal flexibility in estimating their coefficients, particularly with different signs; the number of selected variables is no longer limited by the sample size; and the resulting prediction accuracy is competitive or superior compared to the alternatives. We illustrate the proposed method by extensive simulation studies. The proposed method is also applied to a Glioblastoma microarray data analysis.

1. Introduction. Suppose the training data set consists of n observations $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_i, y_i), \dots, (\mathbf{x}_n, y_n)$, where $\mathbf{x}_i = (x_{i1}, \dots, x_{ip})'$ is a p -dimensional vector of predictors and y_i is the response variable. We consider the following linear model in this article:

$$(1.1) \quad y_i = \beta_1 x_{i1} + \dots + \beta_p x_{ip} + \varepsilon_i,$$

where ε_i is the error term with mean zero. We assume that the response and the predictors are mean-corrected, so we can exclude the intercept term from model (1.1).

Our motivating application comes from the area of microarray data analysis [Horvath et al. (2006)], which embodies some of the properties of the model (1.1) in many modern applications:

Received June 2008; revised June 2010.

¹Supported in part by NSF Grant DMS-07-06700.

²Supported in part by EU Grant MIRG-CT-2007-208019 and ISF Grant 1227/09.

³Supported in part by NSF Grants DMS-07-05532 and DMS-07-48389.

Key words and phrases. Lasso, microarray, regularization, variable selection.

1. In a typical microarray study, the sample size n is usually on the order of 10s, while the number of genes p is on the order of 1000s or even 10,000s. For example, in the glioblastoma microarray gene expression study of Horvath et al. (2006), the sample sizes of the two data sets are 55 and 65, respectively, while the number of genes considered in their analysis is 3600.

2. Microarray data analysis typically combines predictive performance and model interpretation as its goals: one seeks models which explain the phenotype of interest well, but also identify genes, pathways, etc. that might be involved in generating this phenotype.

Shrinkage in general, and variable selection in particular, feature prominently in such applications. Significantly decreasing the number of variables used in the model from the original 1000's to a more manageable number by identifying the most useful and predictive ones usually facilitates both improved accuracy and interpretation.

Variable selection has been studied extensively in the literature; see Breiman (1995), Tibshirani (1996), Fan and Li (2001), Zou and Hastie (2005) and Zou (2006), among many others. In particular, the lasso method proposed by Tibshirani (1996) has gained much attention in recent years.

The lasso criterion penalizes the L_1 -norm of the regression coefficients:

$$(1.2) \quad \min_{\beta} \sum_{i=1}^n \left(y_i - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p |\beta_j|,$$

where λ is a nonnegative tuning parameter. Owing to the singularity of the derivative of L_1 -norm penalty at $\beta_j = 0$, lasso continuously shrinks the estimated coefficients toward zero, and some estimated coefficients will be exactly zero when λ is sufficiently large.

Although lasso has shown success in many situations, it has two limitations in practice [Zou and Hastie (2005)]:

1. When the model includes several highly correlated variables, all of which are related to some extent to the response variable, lasso tends to pick only one or a few of them and shrinks the rest to 0. This may not be a desirable feature. For example, in microarray analysis, expression levels of genes that share one common biological pathway are usually highly correlated, and these genes may all contribute to the biological process, but lasso usually selects only one gene from the group. An ideal method should be able to select all relevant genes, highly correlated or not, while eliminating trivial genes.

2. When $p > n$, lasso can identify at most only n variables before it saturates. This again may not be a desirable feature for many practical problems, particularly microarray studies, for it is unlikely that only such a small number of genes are involved in the development of a complex disease. A method that is able to identify more than n variables should be more desirable for such problems.

Several methods have been proposed recently to alleviate these two possible limitations of lasso mentioned above, including the elastic-net [Zou and Hastie (2005)], the adaptive lasso [Zou (2006)], the relaxed lasso [Meinshausen (2007)] and VISA [Radchenko and James (2008)]. In particular, Zou and Hastie (2005) proposed the elastic-net method, a penalized regression with the mixture of the L_1 -norm and the L_2 -norm penalties of the coefficients:

$$(1.3) \quad \min_{\beta} \sum_{i=1}^n \left(y_i - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda_1 \sum_{j=1}^p |\beta_j| + \lambda_2 \sum_{j=1}^p \beta_j^2,$$

where λ_1 and λ_2 are two nonnegative tuning parameters. Similar to lasso, the elastic-net method also simultaneously does automatic variable selection and continuous shrinkage. Due to the nature of the L_2 -norm penalty, that is, the ridge regression penalty, the number of selected variables is no longer limited by the sample size. However, the ridge penalty forces the estimated coefficients of highly correlated predictors to be close to each other. This feature helps to select or remove highly correlated variables altogether if their coefficients are truly close to each other, but it loses the ability of estimating coefficients of highly correlated variables with different magnitudes, particularly with different signs, which is not rare in practical problems. As a simple illustrative example, eggs are rich in both protein and cholesterol that have quite different effects to human health. When we consider the impact of egg consumption to human health, we have two highly correlated variables with opposite effects. In this scenario, forcing the estimated coefficients of protein and cholesterol to be the same will cause big biases, and is not expected to have adequate prediction performance.

Another modification of the lasso method is the adaptive lasso proposed by Zou (2006), which penalizes the weighted L_1 -norm of the regression coefficients:

$$(1.4) \quad \min_{\beta} \sum_{i=1}^n \left(y_i - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p w_j |\beta_j|,$$

where $w_j = |\hat{\beta}_j^{\text{ols}}|^{-r}$ for a constant $r > 0$, and $\hat{\beta}_j^{\text{ols}}$ is the classical ordinary least squares (OLS) estimator for β_j . Adaptive lasso possesses some nice asymptotic properties that lasso does not have. When p is fixed, n tends to ∞ and λ approaches zero with a certain rate, Zou (2006) has shown that the adaptive lasso approach selects the true underlying model with probability tending to one, and the corresponding estimated coefficients have the same asymptotic normal distribution as they would have if the true underlying model were provided in advance. This is called the “oracle” property by Fan and Li (2001), a property of super-efficiency. Although adaptive lasso has nice asymptotic properties, its finite sample performance does not always dominate lasso because it heavily depends on the precision of the OLS estimation. In his Table 2, Zou [(2006), page 1424] presented that the prediction performance of adaptive lasso can be worse than lasso when the OLS

estimation is more variable. In practice, adaptive lasso suffers (sometimes more severely than lasso) from the multicollinearity caused by large correlations among covariates because OLS estimates are very unstable in this situation. In addition, due to the intrinsic constraint of the L_1 -norm penalty, the number of variables selected by adaptive lasso cannot exceed n .

In this article we propose a novel extension of the lasso method, which we call the random lasso method. The proposed method can handle highly correlated variables in a more flexible manner than elastic-net, especially when their effects have different magnitudes and signs, and also can select more variables than the sample size. Our experiments below demonstrate that the combination of variable selection quality, estimation accuracy, and prediction quality offered by the random lasso is consistently competitive with, and often significantly superior, to those of all the approaches mentioned above. The main price one pays for using the random lasso, however, is in significantly increased computational complexity.

The rest of the paper is organized as follows. We introduce the proposed random lasso method in Section 2, and demonstrate the method using simulation studies in Section 3. In Section 4 we analyze a real data example, and in Section 5 we provide a summary of the proposed method.

2. Random lasso. As mentioned above, one of the limitations of lasso is that it can select only one or a few of a set of highly correlated important variables. If several independent data sets were generated from the same distribution, then we would expect lasso to select nonidentical subsets of those highly correlated important variables from different data sets, and our final collection may be most, or perhaps even all, of those highly correlated important variables by taking a union of selected variables from different data sets. Such a process may yield more than n variables, overcoming the other limitation of lasso.

In practice, however, we have only a single data set at hand, and splitting the available data set into small pieces is not an efficient way of using data. The bootstrap may yield desirable perturbations similar to that of multiple data sets. Because each bootstrap sample may include only a subset of the highly correlated variables, the bootstrap has the ability to break down some of the correlations. Hence, for each bootstrap sample, we can randomly select q candidate variables, with $q \leq p$. This becomes the basic idea of the proposed random lasso approach that has a similar flavor to the random forest method; see [Breiman \(2001\)](#). We also acknowledge that [Park and Hastie \(2008\)](#) proposed using bootstrap to provide a measure of how likely the predictors were to be selected and examine what other predictors could have been preferred. An obvious idea is to build on Park and Hastie's idea to construct a complete predictive modeling tool which may be termed "Bagged Lasso." Our algorithm may be considered a more evolved and "adaptive" version of this idea. In the experiments below we discuss the effects of this added complexity on performance.

Our proposed algorithm is a two-step approach and is described below. In each step, bootstrap samples are drawn to yield the desired perturbation similar to that of multiple data sets. To give the method the most flexibility, we allow different numbers of randomly selected variables to be included in the model, that is, q_1 candidate variables are randomly selected in each bootstrap sample of the first step, and q_2 candidate variables are randomly selected in each bootstrap sample of the second step, where q_1 and q_2 are treated as two tuning parameters that can be chosen as large as p .

ALGORITHM (“Generate” and “Select”). *Step 1.* Generating importance measures for all coefficients:

1a. Draw B bootstrap samples with size n by sampling with replacement from the original training data set.

1b. For the b_1 th bootstrap sample, $b_1 \in \{1, \dots, B\}$, randomly select q_1 candidate variables, and apply lasso to obtain estimators $\hat{\beta}_j^{(b_1)}$ for β_j , $j = 1, \dots, p$. Estimators are zero for coefficients of those unselected variables, either outside the subset of q_1 variables, or excluded by lasso.

1c. Compute the importance measure of x_j by $I_j = |B^{-1} \sum_{b_1=1}^B \hat{\beta}_j^{(b_1)}|$.

Step 2. Selecting variables.

2a. Draw another set of B bootstrap samples with size n by sampling with replacement from the original training data set.

2b. For the b_2 th bootstrap sample, $b_2 \in \{1, \dots, B\}$, randomly select q_2 candidate variables with selection probability of x_j proportional to its importance I_j obtained in step 1c, and apply lasso (or adaptive lasso) to obtain estimators $\hat{\beta}_j^{(b_2)}$ for β_j , $j = 1, \dots, p$. Estimators are zero for coefficients of those unselected variables, either outside the subset of q_2 variables, or excluded by lasso.

2c. Compute the final estimator $\hat{\beta}_j$ of β_j by $\hat{\beta}_j = B^{-1} \sum_{b_2=1}^B \hat{\beta}_j^{(b_2)}$.

In step 1c, we would like to generate an importance score for each predictor to assist variable selection and coefficient estimation in the second step. The average coefficient for each predictor over all bootstrap samples is our choice as an importance score. The intuition is that, for an unimportant variable, the estimated coefficients in different bootstrap samples are likely to be small or even have different signs, so the corresponding average will typically be close to zero. For an important variable, however, the estimated coefficients in different bootstrap samples are likely to be consistently large, and the corresponding average is also large. Therefore, we choose the absolute value of the average as the importance score for each predictor.

In step 2b, there are several choices of weights if adaptive lasso is applied, for example,

$$w_j = 1/|\hat{\beta}_j^{\text{ols}}|^r, \quad w_j = 1/|\hat{\beta}_j^{\text{ridge}}|^r \quad \text{or} \quad w_j = 1/|\hat{\beta}_j^{\text{uni}}|^r,$$

where $\hat{\beta}_j^{\text{ols}}$ is the OLS estimator (if $p < n$), $\hat{\beta}_j^{\text{ridge}}$ is the ridge regression estimator, $\hat{\beta}_j^{\text{uni}}$ is the univariate estimator, and r is a positive number. Instead, we use importance measures obtained in step 1 as the weights for adaptive lasso in our numerical examples and find it works well.

In practice, we need to choose the number of bootstrap samples B , the number of candidate variables to be included in each bootstrap sample q_1 and q_2 , and the tuning parameter λ for (adaptive) lasso to each bootstrap sample. Based on our experience, our algorithm performs similarly when B is large. One can take $B = 500$ or $B = 1000$, for example. We can use cross-validation (CV) to select q_1 and q_2 , and either CV or generalized cross-validation (GCV) to select λ . In the following simulations, we use independent validation data sets.

3. Simulation studies. In this section we use simulations to demonstrate the proposed random lasso method, and compare to a large collection of other methods. Data are generated from model (1.1) with $x_{ij} \sim N(0, 1)$ and $\varepsilon_i \sim N(0, \sigma^2)$.

Five examples are considered. Examples 1 and 2 were used in the lasso paper by Tibshirani (1996), the adaptive lasso paper by Zou (2006), and the elastic-net paper by Zou and Hastie (2005). In Examples 3 and 4, the coefficients of some highly correlated variables have different signs. In Example 5 the number of variables with nonzero coefficients is larger than the sample size. The following are the details of the five examples.

EXAMPLE 1. There are $p = 8$ variables. The pairwise correlation between x_{j_1} and x_{j_2} is set to be $\rho(j_1, j_2) = 0.5^{|j_1 - j_2|}$. We let

$$\beta = (3, 1.5, 0, 0, 2, 0, 0, 0).$$

Following Zou (2006), we consider three values of σ : $\sigma \in \{1, 3, 6\}$. The corresponding signal-to-noise ratios (SNR) are 21.3, 2.4 and 0.6, respectively, where the SNR is defined as $\text{Var}(X'\beta)/\text{Var}(\varepsilon)$.

EXAMPLE 2. We use the same model in Example 1 but with $\beta_j = 0.85$ for all j . We also consider the same three values of σ as in Example 1. The corresponding signal-to-noise ratios (SNR) are 14.5, 1.6 and 0.4, respectively.

EXAMPLE 3. There are $p = 40$ variables. The first 10 coefficients are nonzero. The correlation between each pair of the first 10 variables is set to be 0.9. The remaining 30 variables are independent with each other, and also independent with the first 10 variables. We let

$$\beta = (3, 3, 3, 3, 3, -2, -2, -2, -2, -2, 0, \dots, 0),$$

and $\sigma = 3$. The SNR is about 3.2.

EXAMPLE 4. There are $p = 40$ variables. The first six coefficients are nonzero. The pairwise correlation between the first three variables is set to be 0.9, and the same correlation structure is also set for the second three variables. The remaining 34 variables are independent from each other. The first three variables, the second three variables, and the remaining 34 variables are independent from each other. We let

$$\beta = (3, 3, -2, 3, 3, -2, 0, \dots, 0),$$

and $\sigma = 6$. The SNR is about 0.9.

EXAMPLE 5. There are $p = 120$ variables. The first 60 coefficients are nonzero and drawn from $N(3, 0.5)$, and their values are then fixed for all simulation runs. The remaining 60 coefficients are set to be zero. The covariate matrix is generated from a multivariate normal distribution with zero mean and covariance matrix as

$$\begin{pmatrix} \Sigma_0 & 0 & 0 & 0 \\ 0 & \Sigma_0 & 0.2J & 0 \\ 0 & 0.2J & \Sigma_0 & 0 \\ 0 & 0 & 0 & \Sigma_0 \end{pmatrix},$$

where Σ_0 is a 30×30 matrix with unit diagonal elements and off-diagonal elements of value 0.7, and J is a 30×30 matrix with all unit elements.

For Examples 1–4, we consider two sample sizes: $n = 50$ and $n = 100$. For Example 5, since the purpose is to study the performance of methods under the situation with $p > n$, we consider only sample size $n = 50$. For each example, we also generate a validation data set with the same sample size as the training data set. Models are fitted on training data only, and the validation data are used for selecting the tuning parameters that minimize the prediction error within their context respectively. Regarding the number of bootstrap samples, we used $B = 200$. We also tried $B = 500$; the results are similar to those of $B = 200$.

We calculate the relative model error (RME) given below to evaluate the prediction performance of each predictive model. Suppose that the fitted coefficient vector is $\hat{\beta}$ and the true coefficient vector is β^0 , then the relative model error is defined as follows:

$$\text{Relative Model Error} = (\hat{\beta} - \beta^0)' \Sigma (\hat{\beta} - \beta^0) / \sigma^2,$$

where Σ is the covariance matrix of the predictors, and σ is the standard deviation of the error term in model (1.1).

We repeat the simulation 100 times and compute the average of RMEs and their standard errors. We also record how frequently each variable is selected during the 100 simulations. For the variable selection of random lasso, since the final estimator is the average over all bootstrap samples, it is very easy for a variable to

have a nonzero coefficient if it has a nonzero coefficient in any particular bootstrap sample. So it is a little unfair to use zero or nonzero as the variable selection criterion for random lasso. In this paper we introduce a threshold t_n , and consider a variable x_j to be selected, only if the corresponding coefficient $|\hat{\beta}_j| > t_n$. In the following simulation studies, we chose $t_n = 1/n$, where n is the sample size of the training data. We compare the performance in prediction accuracy and variable selection frequency of random lasso with the following methods: OLS, lasso, adaptive lasso, elastic-net and two other recently developed methods: relaxed lasso [(Meinshausen (2007))] and VISA [Redchenko and James (2008)]. In Example 5, since $n < p$, the OLS estimator is not unique, so we fitted ridge regression, and used the inverse of the absolute value of the ridge regression estimator as the weight for adaptive lasso. Results are summarized in Tables 1 and 2.

As we can see from Table 1, shrinkage methods perform much better than OLS in most cases. This illustrates that some regularization is crucial in achieving higher prediction accuracy. We also see that random lasso has competitive RMEs with all other methods in Examples 1 and 2, except perhaps when compared to elastic-net on Example 2. However, one should keep in mind that Example 2 represents the motivating setup for the elastic-net and, thus, this result is not surprising. Random lasso has consistently smaller RMEs than all other regularization methods in Examples 3–5. It also has the highest important variable selection frequency (see Table 2). In fact, random lasso selects most of the important variables all the time. It also has competitive performance in removing unimportant variables compared to other methods in Examples 1, 3 and 4. In Example 5 random lasso selects more unimportant variables than other methods, but it also selects almost all important variables while other methods perform poorly on this aspect.

It is interesting to compare the elastic net and the random lasso in terms of the signs of the estimated nonzero coefficients of the important variables in Examples 3 and 4. In these two examples, the important variables are highly correlated but with different signs. The result is summarized in Tables 4 and 5. We can see that random lasso has much better performance in estimating correct signs for truly negative coefficients, and much smaller estimation bias than the elastic net method.

For random lasso, the q_1 and q_2 selection can be crucial. For Examples 1 and 2, we select the optimal q_1 and q_2 based on the validation data set among values 2, 4, 6 and 8, for Examples 3 and 4, we select the optimal q_1 and q_2 among values 4, 8, 12, 16, 20, 24 and 28, and for Example 5, we select the optimal q_1 and q_2 among values 5, 10, 15, 20 and 25. We also summarize the frequency for the selected q_1 and q_2 in Examples 1 and 2 with sample size $n = 50$ (see Table 3). From the last columns and the last rows of the six sub-tables, we can see that random lasso prefers a smaller number of predictors in both the first stage and the second stage of the algorithm, as σ becomes larger (correspondingly, the signal-to-noise ratio is smaller). This illustrates that the random subset selection of variables in each bootstrap sample can be helpful, when the signal-to-noise ratio is small.

TABLE 1
1000× Average relative model errors of different methods for all five examples

		OLS	Lasso	ALasso	Enet	Relaxo	VISA	RLasso
<i>Example 1</i>								
<i>n</i> = 50	$\sigma = 1$	212 (13)	131 (9)	92 (8)	132 (9)	88 (7)	91 (8)	82 (7)
	$\sigma = 3$	219 (11)	146 (9)	141 (9)	140 (9)	<i>101</i> (7)	97 (7)	105 (7)
	$\sigma = 6$	201 (11)	119 (6)	131 (7)	<i>104</i> (6)	117 (7)	115 (7)	96 (6)
<i>n</i> = 100	$\sigma = 1$	89 (6)	59 (4)	37 (3)	59 (4)	35 (3)	36 (3)	34 (3)
	$\sigma = 3$	90 (5)	59 (4)	48 (4)	58 (4)	38 (3)	39 (3)	43 (3)
	$\sigma = 6$	89 (5)	56 (5)	59 (5)	<i>49</i> (4)	50 (5)	50 (5)	42 (3)
<i>Example 2</i>								
<i>n</i> = 50	$\sigma = 1$	218 (13)	211 (12)	229 (13)	181 (10)	— —	<i>210</i> (12)	217 (12)
	$\sigma = 3$	202 (10)	171 (8)	200 (10)	140 (8)	— —	180 (9)	<i>167</i> (8)
	$\sigma = 6$	203 (12)	127 (5)	158 (7)	111 (6)	— —	144 (7)	<i>112</i> (4)
<i>n</i> = 100	$\sigma = 1$	82 (4)	84 (4)	90 (4)	77 (4)	— —	<i>81</i> (4)	88 (4)
	$\sigma = 3$	91 (5)	87 (5)	92 (5)	74 (4)	— —	95 (6)	<i>81</i> (7)
	$\sigma = 6$	87 (4)	69 (4)	85 (5)	58 (4)	— —	73 (4)	112 (4)
<i>Example 3</i>								
<i>n</i> = 50		5259 (313)	666 (15)	613 (17)	562 (12)	<i>608</i> (13)	610 (16)	299 (11)
<i>n</i> = 100		680 (20)	505 (11)	<i>313</i> (11)	471 (10)	487 (11)	487 (11)	132 (6)
<i>Example 4</i>								
<i>n</i> = 50		4913 (323)	233 (11)	216 (12)	203 (10)	155 (9)	<i>152</i> (9)	126 (5)
<i>n</i> = 100		706 (25)	144 (6)	122 (5)	115 (5)	100 (5)	98 (5)	70 (36)
<i>Example 5</i>								
<i>n</i> = 50		174 (9)	394 (12)	470 11	241 (11)	395 (11)	421 (12)	227 (11)

Notes: ALasso—the adaptive lasso estimator; Enet—elastic net; RLasso—random lasso. The numbers in the parentheses are the corresponding 1000× standard errors. In each row, we mark the best performing method in bold and the second best in italics.

TABLE 2
Variable selection frequencies (%) of different methods for all five examples

	Lasso	ALasso	Enet	Relaxo	VISA	RLasso
<i>Example 1</i>						
<i>n</i> = 50						
IV ($\sigma = 1$)	(100, 100, 100)	(100, 100, 100)	(100, 100, 100)	(100, 100, 100)	(100, 100, 100)	(100, 100, 100)
UV ($\sigma = 1$)	(46, 58, 64)	(23, 27, 38)	(46, 59, 64)	(10, 15, 19)	(11, 17, 20)	(28, 33, 44)
IV ($\sigma = 3$)	(99, 100, 100)	(95, 99, 100)	(100, 100, 100)	(93, 100, 100)	(97, 100, 100)	(99, 100, 100)
UV ($\sigma = 3$)	(48, 55, 61)	(33, 40, 48)	(44, 55, 69)	(11, 18, 21)	(15, 21, 24)	(45, 57, 68)
IV ($\sigma = 6$)	(76, 85, 99)	(62, 76, 96)	(85, 92, 100)	(60, 70, 98)	(61, 72, 98)	(92, 94, 100)
UV ($\sigma = 6$)	(47, 49, 53)	(32, 36, 38)	(43, 51, 70)	(15, 19, 21)	(15, 19, 24)	(40, 48, 58)
<i>n</i> = 100						
IV ($\sigma = 1$)	(100, 100, 100)	(100, 100, 100)	(100, 100, 100)	(100, 100, 100)	(100, 100, 100)	(100, 100, 100)
UV ($\sigma = 1$)	(54, 59, 64)	(27, 27, 32)	(53, 60, 63)	(13, 16, 25)	(14, 20, 25)	(19, 29, 38)
IV ($\sigma = 3$)	(100, 100, 100)	(100, 100, 100)	(100, 100, 100)	(100, 100, 100)	(100, 100, 100)	(100, 100, 100)
UV ($\sigma = 3$)	(45, 51, 57)	(22, 30, 32)	(44, 55, 67)	(6, 13, 17)	(13, 18, 19)	(36, 47, 56)
IV ($\sigma = 6$)	(96, 99, 100)	(86, 99, 100)	(99, 99, 100)	(90, 93, 100)	(90, 93, 100)	(100, 100, 100)
UV ($\sigma = 6$)	(47, 57, 63)	(36, 40, 47)	(42, 63, 68)	(11, 23, 25)	(11, 25, 27)	(37, 54, 61)
<i>Example 2</i>						
<i>n</i> = 50						
IV ($\sigma = 1$)	(100, 100, 100)	(100, 100, 100)	(100, 100, 100)	–	(100, 100, 100)	(100, 100, 100)
IV ($\sigma = 3$)	(89, 92, 96)	(88, 90, 96)	(92, 96, 99)	–	(83, 88, 90)	(98, 99, 99)
IV ($\sigma = 6$)	(69, 72, 76)	(55, 60, 68)	(72, 78, 88)	–	(50, 54, 65)	(83, 89, 95)
<i>n</i> = 100						
IV ($\sigma = 1$)	(100, 100, 100)	(100, 100, 100)	(100, 100, 100)	–	(100, 100, 100)	(100, 100, 100)
IV ($\sigma = 3$)	(95, 97, 100)	(96, 97, 100)	(98, 99, 100)	–	(92, 95, 97)	(99, 100, 100)
IV ($\sigma = 6$)	(81, 86, 89)	(78, 81, 85)	(72, 78, 88)	–	(50, 55, 65)	(83, 89, 95)

RANDOM LASSO

TABLE 2
(Continued)

	Lasso	ALasso	Enet	Relaxo	VISA	RLasso
<i>Example 3</i>						
<i>n</i> = 50						
IV	(4, 35, 70)	(19, 38, 62)	(20, 60, 95)	(3, 29, 61)	(2, 28, 60)	(93, 98, 100)
UV	(14, 20, 30)	(6, 11, 18)	(6, 13, 18)	(7, 9, 15)	(4, 7, 14)	(10, 17, 24)
<i>n</i> = 100						
IV	(45, 69, 95)	(68, 82, 93)	(51, 76, 99)	(39, 62, 88)	(38, 62, 88)	(98, 99, 99)
UV	(43, 52, 55)	(15, 21, 31)	(29, 35, 40)	(27, 36, 42)	(27, 37, 43)	(22, 30, 37)
<i>Example 4</i>						
<i>n</i> = 50						
IV	(11, 70, 77)	(16, 49, 59)	(63, 92, 96)	(4, 63, 70)	(4, 62, 73)	(84, 96, 97)
UV	(12, 17, 25)	(4, 8, 14)	(9, 17, 23)	(0, 4, 9)	(1, 3, 8)	(11, 21, 30)
<i>n</i> = 100						
IV	(8, 84, 88)	(17, 62, 72)	(70, 98, 99)	(3, 75, 84)	(3, 76, 85)	(89, 99, 99)
UV	(12, 22, 31)	(4, 10, 14)	(7, 14, 21)	(1, 3, 8)	(1, 4, 9)	(8, 14, 21)
<i>Example 5</i>						
IV	(19, 30, 40)	(15, 25, 35)	(40, 50, 61)	(14, 23, 34)	(16, 27, 35)	(76, 86, 95)
UV	(3, 8, 14)	(0, 7, 11)	(1, 5, 8)	(0, 3, 8)	(0, 2, 8)	(18, 29, 38)

Notes: Since OLS always includes all variables, it is excluded from the comparison. IV—important variables; UV—unimportant variables. The three numbers in each pair of parentheses are the min, median, and max of selection frequencies among all important or unimportant variables, respectively.

TABLE 3
Frequencies (%) of the selected q_1 and q_2 for Examples 1 and 2

	$q_2 = 2$	$q_2 = 4$	$q_2 = 6$	$q_2 = 8$	Total
<i>Example 1: $n = 50$</i>					
$\sigma = 1$					
$q_1 = 2$	0	0	3	9	12
$q_1 = 4$	0	2	3	8	13
$q_1 = 6$	0	2	8	9	19
$q_1 = 8$	0	7	21	28	56
Total	0	11	35	54	100
$\sigma = 3$					
$q_1 = 2$	0	9	11	18	38
$q_1 = 4$	0	2	10	10	22
$q_1 = 6$	0	0	7	5	12
$q_1 = 8$	0	1	12	15	28
Total	0	12	40	48	100
$\sigma = 6$					
$q_1 = 2$	8	22	15	17	62
$q_1 = 4$	2	1	5	7	15
$q_1 = 6$	0	1	6	6	13
$q_1 = 8$	0	0	5	5	10
Total	10	24	31	35	100
<i>Example 2: $n = 50$</i>					
$\sigma = 1$					
$q_1 = 2$	0	0	4	39	43
$q_1 = 4$	0	0	1	30	31
$q_1 = 6$	0	0	0	15	15
$q_1 = 8$	0	0	1	10	11
Total	0	0	6	94	100
$\sigma = 3$					
$q_1 = 2$	0	10	24	28	62
$q_1 = 4$	0	1	2	11	14
$q_1 = 6$	0	0	0	12	12
$q_1 = 8$	0	0	1	11	12
Total	0	11	27	62	100
$\sigma = 6$					
$q_1 = 2$	2	23	27	18	70
$q_1 = 4$	0	1	2	5	8
$q_1 = 6$	0	0	4	8	12
$q_1 = 8$	0	0	4	6	10
Total	2	24	37	37	100

It should be noted that we also experimented with “Bagged Lasso” (that is, a 1-step bootstrap approach with $q = p$) on all simulations. The results were reasonable and, in fact, very similar to the elastic net results on all setups. However, since

TABLE 4
Coefficient and coefficient sign estimation of elastic net and random lasso for Example 3

	β_1	β_2	β_3	β_4	β_5	β_6	β_7	β_8	β_9	β_{10}
True coef.	3	3	3	3	3	−2	−2	−2	−2	−2
Enet ($n = 50$)										
Ave. of est.	1.03 (0.07)	1.06 (0.07)	0.91 (0.06)	1.04 (0.08)	0.98 (0.07)	−0.05 (0.06)	−0.03 (0.04)	−0.03 (0.05)	0.01 (0.03)	0.04 (0.02)
Freq. (%) of pos. sgn.	94	91	92	95	91	23	16	17	19	27
Freq. (%) of neg. sgn.	0	0	0	0	0	5	6	3	4	1
RLasso ($n = 50$)										
Ave. of est.	1.84 (0.12)	2.01 (0.12)	1.75 (0.12)	1.81 (0.11)	1.84 (0.11)	−0.84 (0.09)	−0.89 (0.07)	−0.88 (0.07)	−0.91 (0.07)	−0.83 (0.07)
Freq. (%) of pos. sgn.	98	99	96	98	100	9	4	4	7	2
Freq. (%) of neg. sgn.	2	0	2	2	0	88	95	93	93	97
Enet ($n = 100$)										
Ave. of est.	1.42 (0.10)	1.54 (0.09)	1.47 (0.10)	1.43 (0.09)	1.61 (0.11)	−0.53 (0.09)	−0.52 (0.09)	−0.47 (0.09)	−0.38 (0.07)	−0.52 (0.09)
Freq. (%) of pos. sgn.	98	99	98	97	98	15	20	17	17	19
Freq. (%) of neg. sgn.	0	0	0	0	0	37	34	33	34	35
RLasso ($n = 100$)										
Ave. of est.	2.33 (0.09)	2.51 (0.09)	2.45 (0.09)	2.31 (0.08)	2.48 (0.09)	−1.51 (0.07)	−1.35 (0.06)	−1.46 (0.07)	−1.33 (0.06)	−1.41 (0.07)
Freq. (%) of pos. sgn.	99	99	99	98	99	1	0	1	1	0
Freq. (%) of neg. sgn.	0	0	0	0	0	98	99	98	98	99

TABLE 5
Coefficient and coefficient sign estimation of elastic net and random lasso for Example 4

	β_1	β_2	β_3	β_4	β_5	β_6
True coef.	3	3	-2	3	3	-2
Enet ($n = 50$)						
Ave. of est.	1.30 (0.07)	1.44 (0.08)	0.51 (0.06)	1.75 (0.09)	1.47 (0.07)	0.74 (0.07)
No. of pos. sgn.	92	94	63	96	92	70
No. of neg. sgn.	0	0	0	0	0	1
RLasso ($n = 50$)						
Ave. of est.	1.85 (0.12)	1.68 (0.13)	-0.17 (0.07)	2.01 (0.14)	1.89 (0.13)	-0.17 (0.09)
No. of pos. sgn.	98	90	33	91	96	38
No. of neg. sgn.	1	7	65	5	2	57
Enet ($n = 100$)						
Ave. of est.	1.57 (0.06)	1.57 (0.07)	0.54 (0.05)	1.69 (0.06)	1.67 (0.06)	0.61 (0.05)
No. of pos. sgn.	97	98	69	98	99	72
No. of neg. sgn.	0	0	0	0	0	0
RLasso ($n = 100$)						
Ave. of est.	2.25 (0.06)	1.91 (0.07)	-0.57 (0.05)	2.28 (0.06)	2.08 (0.06)	-0.55 (0.05)
No. of pos. sgn.	99	97	17	100	99	15
No. of neg. sgn.	0	2	81	0	0	83

these results are clearly inferior overall to the random lasso, we chose to eliminate them to avoid clutter.

4. Glioblastoma gene expression data analysis. We analyze the data from a glioblastoma microarray gene expression study conducted by Horvath et al. (2006) by using the proposed random lasso method and compare with the lasso, adaptive lasso, relaxed lasso, elastic-net and VISA methods.

Glioblastoma is the most common primary malignant brain tumor of adults and one of the most lethal of all cancers. Patients with this disease have a median survival of 15 months from the time of diagnosis despite surgery, radiation, and chemotherapy. Global gene expression data from two independent sets of clinical tumor samples of $n = 55$ and $n = 65$ are obtained by high-density Affymetrix arrays. Expression values of 3600 genes are available. Among the first set of 55 patients, five were alive at the last followup and four were alive in the second set. In our analysis, we exclude these nine censored subjects, and use the logarithm of time to death as the response variable. The first data set is used as the training set and the second data set as the test set.

We first assess each of the 3600 genes by running simple linear regression on the training set, and then select 1000 genes with the smallest p -values. Starting with these 1000 genes, we fit a linear regression model by the proposed random lasso method on the training set, and select 58 genes. Table 7 lists the gene symbol and estimated coefficient for each of these 58 genes. The model with these selected 58 genes is then used to predict the log-survival times for subjects in the test set. We also analyze the data using other lasso-related methods starting with the same 1000 genes on the training set and evaluate obtained models using the test set.

Table 6 shows the number of genes selected by each of these six methods in the training set and corresponding mean prediction error in the test set. We can see that random lasso has the smallest prediction error. It also selects more genes than the other five methods. Among the 58 genes selected by random lasso, 7 genes are also selected by lasso, adaptive lasso, relaxed lasso, VISA and elastic-net (for adaptive lasso, the adaptive weights were calculated using ridge regression).

Several genes identified by the proposed method are of interest. VSNL1, RGS3 and S100A4 are identified to be negatively associated with the patients' survival. VSNL1 is a member of the visinin/recoverin subfamily of neuronal calcium sensor proteins. The encoded protein is strongly expressed in granule cells of the cerebellum where it associates with membranes in a calcium-dependent manner and modulates intracellular signaling pathways of the central nervous system by directly or indirectly regulating the activity of adenylyl cyclase. A previous study [Xie et al. (2007)] has demonstrated that VSNL1 plays a very important role in neuroblastoma metastasis, and VSNL1 mRNA in highly invasive cells is significantly higher than that in lowly invasive cells. RGS3 encodes a member of the regulator of the G-protein signaling (RGS) family. This protein is a GTP-ase activating protein which inhibits G-protein mediated signal transduction. Tatenhorst et al. (2004) demonstrated that glioma cell clones overexpressing RGS3 showed an increase of both adhesion and migration. S100A4 encodes a member of the S100 family of proteins, which are localized in the cytoplasm and/or nucleus of a wide range of cells, and involved in the regulation of a number of cellular processes such as cell cycle progression and differentiation. It is conjectured that the protein encoded by S100A4 may function in motility, invasion, and metastasis [Zou

TABLE 6
Analysis of the glioblastoma data set

Method	No. of genes selected	Mean prediction error
Lasso	29	1.118 (0.205)
Adaptive lasso	33	1.143 (0.211)
Relaxed lasso	23	1.054 (0.194)
Elastic-net	28	1.113 (0.204)
VISA	15	0.997 (0.188)
Random lasso	58	0.950 (0.210)

TABLE 7

Gene symbol and estimated coefficient for each of the 58 genes selected by random lasso based on 50 subjects in the training set

Gene symbol	Estimated coefficients	Gene symbol	Estimated coefficients
VSNL1	−3.839	KIAA0194	−0.039
SNAP25	−1.561	MOBP	−0.033
UBE2D3	−0.382	PTGDS	−0.028
ARF4	−0.341	KIF5A	−0.024
CSNK1A1	−0.319	GORASP2	−0.021
C13orf11	−0.312	ME2	−0.019
CHGA	−0.310	CGI-141	−0.019
C11orf24	−0.223	p25	−0.017
OPTN	−0.221	UGT8	−0.016
UNC84B	−0.176	CKMT1	−0.014
S100A1	−0.157	KIF1A	−0.013
KCNS1	−0.155	KCNAB2	−0.012
NPY	−0.124	C3orf4	−0.011
TIP-1	−0.107	DNASE1L1	−0.011
FAIM2	−0.086	RNF44	0.011
FSTL3	−0.074	ATP6V1B2	0.012
NEFH	−0.072	POLR3E	0.012
CTSK	−0.071	LIN7C	0.014
RGS3	−0.071	GBP2	0.015
PGCP	−0.070	CSF1R	0.018
FLJ20254	−0.059	JIK	0.019
ANXA2	−0.053	—	0.019
FLJ11155	−0.052	C1S	0.026
P2RX4	−0.049	ARHGAP15	0.040
GNPMB	−0.044	PPM1H	0.063
ICAM5	−0.043	MARK4	0.071
ADIPOR1	−0.043	HPCAL4	0.196
BSCL2	−0.042	SULT4A1	0.785
AMBP	−0.042	BSN	2.662

et al. (2005)]. VSNL1, RGS3 and S100A4 were also identified to be related to the poor survival of brain tumor patients in Freije et al. (2004). BSN is identified to be positively associated with the patients' survival. This gene is expressed primarily in neurons in the brain, and the protein encoded by this gene is thought to be a scaffolding protein involved in organizing the presynaptic cytoskeleton. Additional studies will be required to establish the direct relationships between the expression of these genes and the Glioblastoma tumor behavior.

It is also interesting to observe that estimated coefficients of VSNL1 and BSN (−3.839 and 2.662, resp.) have different signs, but the correlation between the expression levels for these two genes in the training set is very high ($\rho = 0.96$). Neither lasso nor elastic-net picked up these two genes. It is worth conducting

more detailed experiments to further explore the connection between VSNL1 and BSN, and their relations to the Glioblastoma tumor behavior.

5. Conclusion. We have proposed the random lasso method for variable selection. The idea of random lasso is mimicking the random forest method [Breiman (2001)] for linear regression models. By drawing bootstrap samples from the original training set and randomly selecting candidate variables, the average of the predictive models based on multiple bootstrap samples alleviates two possible limitations of lasso. It tends to select or remove highly correlated variables more efficiently and has more flexibility in estimating their coefficients than the elastic-net method. The number of variables selected by random lasso is not limited by the sample size. Simulation studies show that the proposed random lasso method has good prediction performance compared to a large set of competitor approaches, and the analysis of Glioblastoma microarray data set demonstrates the usefulness of the proposed method in practice.

Acknowledgments. We would like to thank the Editor, an Associate Editor and three reviewers for their thoughtful and useful comments.

REFERENCES

- BREIMAN, L. (1995). Better subset regression using the non-negative garrote. *Technometrics* **37** 373–384. [MR1365720](#)
- BREIMAN, L. (2001). Random forest. *Mach. Learn.* **45** 5–32.
- FAN, J. and LI, R. (2001). Variable selection via nonconcave penalized likelihood and its oracle properties. *J. Amer. Statist. Assoc.* **96** 1348–1360. [MR1946581](#)
- FREIJE, W., CASTRO-VARGAS, F. E., FANG, Z., HORVATH, S., CLOUGHESY, T., LIAU, L. M., MISCHER, P. S. and NELSON, S. F. (2004). Gene expression profiling of Gliomas strongly predicts survival. *Cancer Res.* **64** 6503–6510.
- HORVATH, S., ZHANG, B., CARLSON, M., LU, K. V., ZHU, S., FELCIANO, R. M., LAURANCE, M. F., ZHAO, W., SHU, Q., LEE, Y., SCHECK, A. C., LIAU, L. M., WU, H., GESCHWIND, D. H., FEBBO, P. G., KORNBLUM, H. I., CLOUGHESY, T. F., NELSON, S. F. and MISCHER, P. S. (2006). Analysis of oncogenic signaling networks in glioblastoma identifies aspm as a novel molecular target. *Proc. Natl. Acad. Sci.* **103** 17402–17407.
- MEINSHAUSEN, N. (2007). Relaxed lasso. *Comput. Statist. Data Anal.* **52** 374–393. [MR2409990](#)
- PARK, M. and HASTIE, T. (2008). Penalized logistic regression for detecting gene interactions. *Biostatistics* **9** 30–50.
- RADCHENKO, P. and JAMES, G. (2008). Variable inclusion and Shrinkage algorithms. *J. Amer. Statist. Assoc.* **103** 1304–1315. [MR2462899](#)
- TATENHORST, L., SENNER, V., PUTTMANN, S. and PAULUS, W. (2004). Regulators of G-protein signaling 3 and 4 (RGS3, RGS4) are associated with Glioma cell motility. *J. Neuropathol. Exp. Neurol.* **63** 210–222.
- TIBSHIRANI, R. (1996). Regression shrinkage and selection via the lasso. *J. Roy. Statist. Soc. Ser. B* **58** 267–288. [MR1379242](#)
- XIE, Y., CHAN, H., FAN, J., CHEN, Y., YOUNG, J., LI, W., MIAO, X., YUAN, Z., WANG, H., TAM, P. K. and REN, Y. (2007). Involvement of visinin-like protein-1 (VSNL-1) in regulating proliferative and invasive properties of neuroblastoma. *Carcinogenesis* **28** 2122–2130.

- ZOU, H. (2006). The adaptive lasso and its oracle properties. *J. Amer. Statist. Assoc.* **101** 1418–1429. [MR2279469](#)
- ZOU, M., AL-BARADIE, R. S., AL-HINDI, H., FARID, N. R. and SHI, Y. (2005). Gene overexpression is associated with invasion and metastasis of papillary thyroid carcinoma. *Br. J. Cancer* **93** 1277–1284.
- ZOU, H. and HASTIE, T. (2005). Regularization and variable selection via the elastic net. *J. Roy. Statist. Soc. Ser. B* **67** 301–320. [MR2137327](#)

S. WANG
DEPARTMENT OF BIOSTATISTICS
UNIVERSITY OF WISCONSIN
MADISON, WISCONSIN, 53792
USA
E-MAIL: swang@biostat.wisc.edu

S. ROSSET
SCHOOL OF MATHEMATICAL SCIENCES
TEL AVIV UNIVERSITY
TEL AVIV
ISRAEL
E-MAIL: saharon@post.tau.ac.il

B. NAN
DEPARTMENT OF BIOSTATISTICS
UNIVERSITY OF MICHIGAN
ANN ARBOR, MICHIGAN, 48109
USA
E-MAIL: bnan@umich.edu

J. ZHU
DEPARTMENT OF STATISTICS
UNIVERSITY OF MICHIGAN
ANN ARBOR, MICHIGAN, 48109
USA
E-MAIL: jizhu@umich.edu